

Intro. Comp. for Data Science (FMI08)

Dr. Nono Saha

May 24, 2023

Max Planck Institute for Mathematics in the Sciences
University of Leipzig/ScaDS.AI

Spring 2023

1. Introduction to **pandas**
2. Series
3. DataFrame
4. Missing value problems and native NAs

Introdcution to pandas

pandas?

pandas implements data frames in **Python** - it takes much of its inspiration from R and NumPy.

pandas aims to be the fundamental high-level building block for practical, real-world data analysis in **Python**. Additionally, it seeks to become the most powerful and flexible open-source data analysis/manipulation tool available in any language.

Key features:

- DataFrame object class
- Reading and writing tabular data
- Data munging (filtering, grouping, summarizing, joining, etc.)
- Data reshaping

pandas - Series

Series - create a pandas series from a list

The columns of a DataFrame are constructed as **Series** - a 1d array-like object containing values of the same type (similar to an ndarray).

```
1 import pandas as pd
```

```
1 pd.Series([1,2,3,4])
```

```
2 ## 0      1
```

```
3 ## 1      2
```

```
4 ## 2      3
```

```
5 ## 3      4
```

```
6 ## dtype: int64
```

```
7
```

```
8
```

```
9 pd.Series(["C","B","A"])
```

```
10 ## 0      C
```

```
11 ## 1      B
```

```
12 ## 2      A
```

```
13 ## dtype: object
```

```
pd.Series(range(5))
```

```
## 0      0
```

```
## 1      1
```

```
## 2      2
```

```
## 3      3
```

```
## 4      4
```

```
## dtype: int64
```

```
pd.Series([1,"A",True])
```

```
## 0      1
```

```
## 1      A
```

```
## 2     True
```

```
## dtype: object
```

pandas - Series: basic methods

Once constructed the components of a series can be accessed via `array()` and `index()` methods.

```
1 s = pd.Series([4,2,1,3])
2 s.array
3 ## <PandasArray>
4 ## [4, 2, 1, 3]
5 ## Length: 4, dtype: int64
6
7 s.index
8 ## RangeIndex(start=0, stop=4, step=1)
```

An index can also be explicitly provided when constructing a Series,

```
1 t = pd.Series([4,2,1,3], index=[
2     "a", "b", "c", "d"])
3 ## a      4
4 ## b      2
5 ## c      1
6 ## d      3
7 ## dtype: int64

1 t.array
2 ## <PandasArray>
3 ## [4, 2, 1, 3]
4 ## Length: 4, dtype: int64
5
6 t.index
```

pandas - Series + NumPy

Series objects are compatible with **NumPy** like functions (vectorized)

```
1 t = pd.Series([4,2,1,3], index=["a","b","c","d"])
```

```
1 t+1
2 ## a      5
3 ## b      3
4 ## c      2
5 ## d      4
6 ## dtype: int64
7
8 np.log(t)
9 ## a      1.386294
10 ## b      0.693147
11 ## c      0.000000
12 ## d      1.098612
13 ## dtype: float64
```

```
1 np.exp(-t**2/2)
2 ## a      0.000335
3 ## b      0.135335
4 ## c      0.606531
5 ## d      0.011109
6 ## dtype: float64
```


pandas - Series indexing

Series can be indexed in the same way as **NumPy** arrays with the addition of being able to use label(s) when selecting elements.

```
1 t = pd.Series([4,2,1,3], index=["a","b","c","d"])
```

```
1 t[1]
2 ## 2
3
4 t[[1,2]]
5 ## b    2
6 ## c    1
7 ## dtype: int64
8
9 t["c"]
10 ## 1
11
12 t[["a","d"]]
13 ## a    4
14 ## d    3
15 ## dtype: int64
```

```
1 t[t == 3]
2 ## d    3
3 ## dtype: int64
4
5 t[t % 2 == 0]
6 ## a    4
7 ## b    2
8 ## dtype: int64
9
10 t["d"] = 6
11 t
12 ## a    4
13 ## b    2
14 ## c    1
15 ## d    6
```

pandas - Series: index alignment

When performing (arithmetic) operations on Series, they will attempt to align by their index. Let us consider the following **Series**:

```
1 m = pd.Series([1,2,3,4], index = ["a","b","c","d"])
2 n = pd.Series([4,3,2,1], index = ["d","c","b","a"])
3 o = pd.Series([1,1,1,1,1], index = ["b","d","a","c","e"])
```

Questions

For each of the following operations, what is the output?

- $m + n$
- $n + m$
- $n + 0$
- $m + 0$

pandas - Series and dicts

We can construct Series from **dicts**, in which case the keys are used to form the index,

```
1 d = {"anna": "A+", "bob": "B-", "carol": "C", "dave": "D+"}
2 pd.Series(d)
3 ## anna      A+
4 ## bob       B-
5 ## carol     C
6 ## dave      D+
7 ## dtype: object
```

Index order will follow key order unless overridden by index,

```
1 pd.Series(d, index = ["dave", "carol", "bob", "anna"])
2
3 ## dave      D+
4 ## carol     C
5 ## bob       B-
6 ## anna      A+
7 ## dtype: object
```

pandas - Series: string series

Series containing strings can be accessed via the `str` attribute,

```
1 s = pd.Series(["the quick", "
    brown fox", "jumps over", "a
    lazy dog"])
2 s
3 ## 0      the quick
4 ## 1      brown fox
5 ## 2      jumps over
6 ## 3      a lazy dog
7 ## dtype: object
8
9 s.str.upper()
10 ## 0      THE QUICK
11 ## 1      BROWN FOX
12 ## 2      JUMPS OVER
13 ## 3      A LAZY DOG
14 ## dtype: object
```

```
1 s.str.split(" ").str[1]
2 ## 0      quick
3 ## 1      fox
4 ## 2      over
5 ## 3      lazy
6 ## dtype: object
7
8 pd.Series([1,2,3]).str
9 ## AttributeError: Can only use
    .str accessor with string
```

pandas - Series: categorical Series

```
1 pd.Series(["Mon",  
            "Tue", "Wed",  
            , "Thur", "  
            Fri"]])  
2 ## 0      Mon  
3 ## 1      Tue  
4 ## 2      Wed  
5 ## 3      Thur  
6 ## 4       Fri  
7 ## dtype: object
```

```
pd.Series(["Mon", "Tue", "Wed", "Thur", "Fri"],  
          dtype="category")  
## 0      Mon  
## 1      Tue  
## 2      Wed  
## 3      Thur  
## 4       Fri  
## dtype: category  
## Categories (5, object): ['Fri', 'Mon', 'Thur',  
                           ', 'Tue', 'Wed']
```

```
1 pd.Series(["Mon", "Tue", "Wed", "Thur", "Fri"], dtype=pd.  
            CategoricalDtype(ordered=True))  
2 ## 0      Mon  
3 ## 1      Tue  
4 ## 2      Wed  
5 ## 3      Thur  
6 ## 4       Fri  
7 ## dtype: category  
8 ## Categories (5, object): ['Fri' < 'Mon' < 'Thur' < 'Tue' < '  
                           Wed']
```

pandas - DataFrame

pandas - DataFrame: constructing DataFrames

panda data frames can also be constructed via `DataFrame()`,
general this is done via dict of columns:

```
1 n = 5
2 d = {
3     "id":      np.random.randint(100, 999, n),
4     "weight":  np.random.normal(70, 20, n),
5     "height":  np.random.normal(170, 15, n),
6     "date":    pd.date_range(start='2/1/2022', periods=n, freq='D')
7 }
8 df = pd.DataFrame(d)
9 df
10
11 ##      id      weight      height      date
12 ## 0   168  102.915535  188.677769  2022-02-01
13 ## 1   615   71.767364  155.907801  2022-02-02
14 ## 2   346   76.666059  171.386839  2022-02-03
15 ## 3   390   74.735465  173.151008  2022-02-04
16 ## 4   556   50.538488  183.083407  2022-02-05
```

pandas - DataFrame: from ndarray

For 2d ndarrays, it is also possible to construct a DataFrame - generally, providing column names and row names (indexes) is a good idea.

```
1 pd.DataFrame(  
2 np.diag([1,2,3]),  
3 columns = ["x","y","z"]  
4 )
```

```
5 ##      x  y  z  
6 ## 0    1  0  0  
7 ## 1    0  2  0  
8 ## 2    0  0  3
```

```
10 pd.DataFrame(  
11 np.diag([1,2,3]),  
12 columns = ["x","y","z"]  
13 )
```

```
14 ##      x  y  z  
15 ## 0    1  0  0  
16 ## 1    0  2  0  
17 ## 2    0  0  3
```

```
1 pd.DataFrame(  
2 np.tri(5,3,-1),  
3 columns = ["x","y","z"],  
4 index = ["a","b","c","d","e"]  
5 )
```

```
6 ##          x      y      z  
7 ## a    0.0    0.0    0.0  
8 ## b    1.0    0.0    0.0  
9 ## c    1.0    1.0    0.0  
10 ## d    1.0    1.0    1.0  
11 ## e    1.0    1.0    1.0
```


pandas - DataFrame: indexing

```
1 df[0]
2 ## KeyError: 0
3
4 df["id"]
5 ## 0      168
6 ## 1      615
7 ## 2      346
8 ## 3      390
9 ## 4      556
10 ## Name: id, dtype: int64
11
12 df.id
13 ## 0      168
14 ## 1      615
15 ## 2      346
16 ## 3      390
17 ## 4      556
18 ## Name: id, dtype: int64
```

```
df[1:3]
##      id      weight      height
      date
## 1   615   71.767364   155.907801
      2022-02-02
## 2   346   76.666059   171.386839
      2022-02-03

df[0::2]
##      id      weight      height
      date
## 0   168   102.915535   188.677769
      2022-02-01
## 2   346    76.666059   171.386839
      2022-02-03
## 4   556    50.538488   183.083407
      2022-02-05
```

pandas - DataFrame: indexing by position

```
1 df.iloc[1] && df.iloc[[1]]
2 ## What is the difference between
   the two instructions?
3
4 df.iloc[0:2]
5 ## id  weight  height  date
6 ## 0   168   102.915535  188.677769
   2022-02-01
7 ## 1   615    71.767364  155.907801
   2022-02-02
8
9 df.iloc[lambda x: x.index % 2 != 0]
10 ##      id      weight      height
      date
11 ## 1   615    71.767364  155.907801
   2022-02-02
12 ## 3   390    74.735465  173.151008
   2022-02-04
```

```
df.iloc[1:3,1:3]
##      weight      height
## 1   71.767364  155.907801
## 2    76.666059  171.386839

df.iloc[0:3, [0,3]]
##      id      date
## 0   168  2022-02-01
## 1   615  2022-02-02
## 2   346  2022-02-03

df.iloc[0:3, [True, True,
               False, False]]
##      id      weight
## 0   168   102.915535
## 1   615    71.767364
## 2   346    76.666059
```

pandas - DataFrame: indexing by name

```
1 df.loc[["anna"]]
2 ## id weight height date
3 ## anna 168 102.915535
   188.677769 2022-02-01
4
5 df.loc["bob":"dave"]
6 ## id weight height date
7 ## bob 615 71.76 155.9 2022-02-02
8 ## carol 346 76.6 171.3 2022-02-03
9 ## dave 390 74.73 173.15 2022-02-04
10
11 df.loc[df.id < 300]
12 ## id weight height date
13 ## anna 168 102.91 188.67
   2022-02-01
```

```
df.loc[:, "date"]
## anna 2022-02-01
## bob 2022-02-02
## carol 2022-02-03
## dave 2022-02-04
## erin 2022-02-05
## Name: date, dtype:
   datetime64[ns]

df.loc[["bob", "erin"], "
   weight":"height"]
## weight height
## bob 71.76 155.9
## erin 50.53 183.08

df.loc[0:2, "weight":"height"
   ]
## ???
```

pandas - DataFrame: views and copies

In general, most pandas operations will generate a new object, but some will return views, mostly the later occurring with subsetting.

```
1 d = pd.DataFrame(np.arange(6).reshape(3,2), columns = ["x", "y"])
```

```
1 v = d.iloc[0:2,0:2]
```

```
2 v
```

```
3 ##      x  y
```

```
4 ## 0    0  1
```

```
5 ## 1    2  3
```

```
6
```

```
7 d.iloc[0,1] = -1
```

```
8 v
```

```
9 ##      x  y
```

```
10 ## 0    0 -1
```

```
11 ## 1    2  3
```

```
v.iloc[0,0] = np.pi
```

```
v
```

```
##              x  y
```

```
## 0  3.141593 -1
```

```
## 1  2.000000  3
```

```
d
```

```
##      x  y
```

```
## 0    0 -1
```

```
## 1    2  3
```

```
## 2    4  5
```

pandas - DataFrame: filtering rows

The `query()` method can be used for filtering rows. It evaluates a string expression in the context of the data frame.

```
1 df = pd.DataFrame(d)
2 df.query('date == "2022-02-01"')
3 ##          id          weight          height          date
4 ## anna    168    102.915535    188.677769  2022-02-01
5
6 df.query('weight > 50')
7 ##          id          weight          height          date
8 ## anna    168    102.915535    188.677769  2022-02-01
9 ## bob     615     71.767364    155.907801  2022-02-02
10 ## carol   346     76.666059    171.386839  2022-02-03
11 ## dave    390     74.735465    173.151008  2022-02-04
12
13 df.query('weight > 50 & height < 165')
14 ##          id          weight          height          date
15 ## bob     615     71.767364    155.907801  2022-02-02
16
17 qid = 414
18 df.query('id == @qid')
```

pandas - DataFrame: element access

```
1 df
2 ##           id           weight           height           date
3 ## anna      168      102.915535      188.677769  2022-02-01
4 ## bob       615       71.767364      155.907801  2022-02-02
5 ## carol     346       76.666059      171.386839  2022-02-03
6 ## dave      390       74.735465      173.151008  2022-02-04
7 ## erin      556       50.538488      183.083407  2022-02-05
```

```
1 df[0,0]
2 ## KeyError: (0, 0)
```

```
4 df.iat[0,0]
5 ## 168
```

```
7 df.id[0]
8 ## 168
```

```
10 df[0:1].id[0]
11 ## 168
```

```
1 df["anna", "id"]
2 ## KeyError: ('anna', 'id')
```

```
4 df.at["anna", "id"]
5 ## 168
```

```
7 df["id"]["anna"]
8 ## 168
```

```
10 df["id"][0]
11 ## 168
```

pandas - DataFrame: properties of a DataFrame

```
1 df
2 ##           id           weight           height           date
3 ##  anna      168      102.915535      188.677769  2022-02-01
4 ##  bob       615       71.767364      155.907801  2022-02-02
5 ##  carol     346       76.666059      171.386839  2022-02-03
6 ##  dave      390       74.735465      173.151008  2022-02-04
7 ##  erin      556       50.538488      183.083407  2022-02-05
```

```
1 df.size
2 ## 20
3
4 df.shape
5 ## (5, 4)
6
7 df.info()
8 ## Try it on your computer and
   analyze the output.
```

```
1 df.dtypes
2 ## id                                int64
3 ## weight                            float64
4 ## height                            float64
5 ## date                             datetime64[ns]
6 ## dtype: object
7
8 df.describe()
```

pandas - DataFrame: selecting columns

Beyond the use of `loc()` and `iloc()`, there is also the `filter()` method which can be used to select columns (or indices) by name with pattern-matching.

```
1 df.filter(items=["id", "weight"])
2 ##           id           weight
3 ## anna      168    102.915535
4 ## bob       615     71.767364
5 ## carol     346     76.666059
6 ## dave      390     74.735465
7 ## erin      556     50.538488
8
9 df.filter(like = "i")
10 ## id weight height
11 ## anna    168   102.91  188.67
12 ## bob     615    71.76  155.90
13 ## carol   346    76.66  171.38
14 ## dave    390    74.73  173.15
15 ## erin    556    50.53  183.08
```

```
1 df.filter(regex="ght$")
2 ## weight height
3 ## anna  102.91  188.67
4 ## bob   71.76  155.90
5 ## carol  76.66  171.38
6 ## dave  74.73  173.15
7 ## erin  50.53  183.08
8
9 df.filter(like="o", axis=0)
10 ## id weight height date
11 ## bob  615  71.76  155.90
12      2022-02-02
13 ## carol 346  76.66  171.38
14      2022-02-03
```


pandas - DataFrame: adding columns

Indexing with assignment allows for in-place modification of a DataFrame, while `assign()` creates a new object (but is chainable).

```
1 df['student'] = [True, True, True,
2                 False, None]
3 df['age'] = [19, 22, 25, None, None]
4 df
```

```
5 ## id weight height date student
6 ##      age
7 ## anna   168   102.91  188.67
8 ##      2022-02-01      True   19.0
9 ## dave   390    74.7   173.15
10 ##      2022-02-04     False   NaN
11 ## erin   556    50.5   183.08
12 ##      2022-02-05     None   NaN
```

```
df.assign(student = lambda x:
            np.where(x.student, "yes", "no"),
            rand = np.random.rand(5)
            )
```

```
## id weight      height      date
##      student age and
## anna   168   102.91  188.67
##      2022-02-01      yes  19.0
##      0.60
## bob    615    71.76  155.90
##      2022-02-02      yes  22.0
##      0.54
```

pandas - DataFrame: removing columns (and rows)

Columns can be dropped via the `drop()` method,

```
1 df.drop(['student'])
2 ## KeyError: "['student'] not
   found in axis."
3
4 df.drop(['student'], axis=1)
5 ## id weight height date age
6 ## anna    168   102.91 188.67
   2022-02-01   19.0
7 ## bob     615    71.76 155.90
   2022-02-02   22.0
8 ## carol   346    76.66 171.38
   2022-02-03   25.0
9 ## dave    390    74.73 173.15
   2022-02-04   NaN
10 ## erin    556    50.53 183.08
   2022-02-05   NaN
```

```
df.drop(columns = df.columns ==
         "age")
## KeyError: '[False, False,
   False, False, False, True]
   not found in axis'

df.drop(columns = df.columns[df.
   columns == "age"])
## id weight height date
   student
## anna    168   102.9 188.67
   2022-02-01    True
## bob     615    71.7 155.9
   2022-02-02    True
## carol   346    76.6 171.3
   2022-02-03    True
```

pandas - DataFrame: row binds

DataFrames can have their rows joined via the `concat()` function (`append()` is also available but deprecated),

```
1 df1 = pd.DataFrame(np.arange(6).  
    reshape(3,2), columns=list("xy"))  
2 df1  
3 ##      x  y  
4 ## 0    0  1  
5 ## 1    2  3  
6 ## 2    4  5  
7  
8 pd.concat([df1,df2])  
9 ##      x  y  
10 ## 0    0  1  
11 ## 1    2  3  
12 ## 2    4  5  
13 ## 0   12 11  
14 ## 1   10  9  
15 ## 2    8  7
```

```
df2 = pd.DataFrame(np.arange  
    (12,6,-1).reshape(3,2),  
    columns=list("xy"))
```

df2

```
##      x  y  
## 0   12 11  
## 1   10  9  
## 2    8  7
```

```
pd.concat([df1.loc[:,["y","x"]],  
    df2])
```

```
##      y  x  
## 0    1  0  
## 1    3  2  
## 2    5  4  
## 0   11 12  
## 1    9 10
```

Missing value problems and native Nas

pandas - Series: missing values

pandas encodes missing values using **NaN** (mostly),

```
1 s = pd.Series({"anna": "A+", "
2               bob": "B-", "carol": "C", "
3               dave": "D+"},
4               index = ["erin", "dave", "carol", "
5                       bob", "anna"])
6
7 s
8 ## erin      NaN
9 ## dave      D+
10 ## carol     C
11 ## bob       B-
12 ## anna      A+
13 ## dtype: object
14
15 pd.isna(s) = ?
```

```
1 s = pd.Series({"anna": 97, "bob"
2               : 82, "carol": 75, "dave":
3               68},
4               index = ["erin", "dave", "carol", "
5                       bob", "anna"], dtype='int64')
6
7 s
8 ## erin      NaN
9 ## dave      68.0
10 ## carol     75.0
11 ## bob       82.0
12 ## anna      97.0
13 ## dtype: float64
14
15 pd.isna(s) = ?
```

pandas - Series: aside - why np.isna()?

```
1 s = pd.Series([1,2,3,None])
2 s
3 ## 0      1.0
4 ## 1      2.0
5 ## 2      3.0
6 ## 3      NaN
7 ## dtype: float64
```

```
1 pd.isna(s)
2 ## 0      False
3 ## 1      False
4 ## 2      False
5 ## 3       True
6 ## dtype: bool
7
8 s == np.nan
9 ## 0      False
10 ## 1      False
11 ## 2      False
12 ## 3      False
13 ## dtype: bool
14
```

```
1 np.nan == np.nan
2 ## What about this?
3
4 np.nan != np.nan
5 ## And this?
6
```

pandas - Series: native NAs

Recent versions of pandas have attempted to adopt a more native missing value, particularly for integer and boolean types,

```
1 pd.Series([1,2,3,None])
2 ## 0      1.0
3 ## 1      2.0
4 ## 2      3.0
5 ## 3      NaN
6 ## dtype: float64
7
8 pd.Series([True,False,None])
9 ## 0      True
10 ## 1     False
11 ## 2      None
12 ## dtype: object
13
```

```
1 pd.isna( pd.Series([1,2,3,None])
2         )
3 ## 0     False
4 ## 1     False
5 ## 2     False
6 ## 3      True
7 ## dtype: bool
8
9 pd.isna( pd.Series([True,False,
10                    None]))
11 ## 0     False
12 ## 1     False
13 ## 2      True
14 ## dtype: bool
```